Bayesian Inference II

Introduction to Stan

Leonardo Egidi, legidi@units.it

2025/2026

Guest lecture, February 11

Table of contents i

Stanislaw Ulam (1909-1984): Manhattan project, H-Bomb experiments in Los Alamos, MCMC father jointly with John von Neumann.

What is Stan?

What is Stan?

- Probabilistic programming language and inference algorithms.
- Stan program
  - declares data and (constrained) parameter variables
  - defines log posterior (or penalized likelihood)
- Stan inference
  - MCMC for full Bayes (`rstan`, `rstanarm`, `brms`, `pyStan`)
  - Variational Bayes for approximate Bayes (`cmdstan`)
  - Optimization for (penalized) MLE (`cmdstan`)
- Stan ecosystem
  - lang, math library (C++)
  - interfaces and tools (R, Python, Julia, many more)
  - documentation (example model repo, ▶ user guide & reference manual ,   , R package vignettes)
 - online community (▶ Stan Forums on Discourse)

4

Why Stan?

- Fit rich Bayesian statistical models. Close to the *big data* philosophy.
- Efficiency
    - Hamiltonian Monte Carlo + NUTS
    - Compiled to C++
- Flexible domain specific language
- "Freedom-respecting, open-source"
    - doc & written materials
    - interacting community
    - continuous development
- Interaction with some other R packages designed to explore the Stan output.

Who is using Stan?

- Biological & physical sciences: clinical trials, epidemiology, genomics, population ecology, entomology, ophthalmology, neurology, agriculture, fisheries, cancer biology, astrophysics & cosmology, molecular biology, oceanography, climatology.

- Social sciences: population dynamics. psycholinguistics, social networks, political science, human development, economics.

- Many more: sports analytics, public health, publishing, finance, pharma, actuarial, recommender systems, educational testing, materials engineering.

Why Stan?

With Stan, we aim to provide an MCMC implementation that works robustly for as many target distributions as possible

- Gibbs, RW Metropoilis can be very inefficient, hard to diagnose.

- To explore complicated high-dimensional spaces we need to leverage what we know about the geometry of the typical set.

- For such a reason, Stan enjoys **Hamiltonian Monte Carlo**.

We will have now only a brief sketch about how HMC works (further readings are mentioned later on). The Stan users, however, may use, analyze and interpret HMC outputs as they were *standard* MCMC outputs!

Moving to Hamiltonian Monte Carlo

Once we have built a model, Bayesian computation reduces to evaluating expectations, or integrals.

$$E_\pi(\theta|y) = \int \theta \pi(\theta|y)d\theta \qquad (1)$$

How do we compute posterior expectations in practice?

- Construct a Markov chain that explores the parameter space.

- Anything you would want to do if you could write it analytically, you can do to any accuracy with the draws (history) of the chain

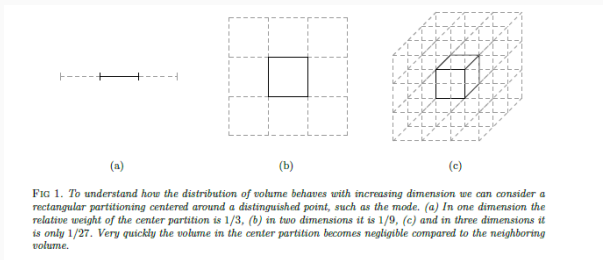$$\lim_{S \to \infty} \frac{1}{S} \sum_{s=1}^{S} \theta^{(s)} \to E_\pi(\theta|y)$$

To be efficient we need to focus computation on the relevant neighborhoods of parameter space. Relevant neighborhoods, however, are defined not by probability density but rather by probability mass.

But exactly which neighborhoods end up contributing most to arbitrary expectations?

The neighborhoods around the maxima of probability distributions feature a lot of probability density, but, especially in a large number of dimensions, or in long tailed distributions, they do not feature much volume. In other words, the *sliver* size $d\theta$ tends to be small there.

The Geometry of High-Dimensional Spaces

Consider a rectangular partitioning centered around a distinguished point, such as the mode (example from Betancourt, 2017):

FIG 1. *To understand how the distribution of volume behaves with increasing dimension we can consider a rectangular partitioning centered around a distinguished point, such as the mode. (a) In one dimension the relative weight of the center partition is 1/3, (b) in two dimensions it is 1/9, (c) and in three dimensions it is only 1/27. Very quickly the volume in the center partition becomes negligible compared to the neighboring volume.*

One of the characteristic properties of high-dimensional spaces is that there is much more volume outside any given neighborhood than inside of it!

Thus, relevant neighborhoods are defined not by probability density but rather by probability mass.

Probability mass concentrates on a hypersurface called the typical set that surrounds the mode.

Moving to Hamiltonian Monte Carlo

To accurately estimate expectations we need a method for numerically finding and then exploring the typical set.

A Markov transition that targets our desired distribution naturally concentrates towards probability mass. An inherent inefficiency in the Gibbs sampler and in the random walk Metropolis Hastings is their random walk behaviour: the simulations can take a long time *zigging* and *zagging* while moving through the target distribution.

HMC borrows strengths from physics to suppress the random walk behaviour in the Metropolis algorithm, thus allowing it to move much more rapidly through the target distribution. The method enjoys the gradient of the log-posterior distribution, $\frac{d \log(\pi(\theta|y))}{d\theta}$ for a sort of adjustment of the algorithm towards the typical set area.

Under ideal conditions, MCMC estimators converge to the true expectations in a very practical progression.

Under ideal conditions, MCMC estimators converge to the true expectations in a very practical progression.

Under ideal conditions, MCMC estimators converge to the true expectations in a very practical progression.

There are many pathological posterior geometries, however, that spoil these ideal conditions.

There are many pathological posterior geometries, however, that spoil these ideal conditions.

There are many pathological posterior geometries, however, that spoil these ideal conditions.

- Hamiltonian Monte Carlo yields fast, and *robust*, exploration of the distributions common in practice.

Comparison of algorithms on **highly correlated** 250-dimensional Gaussian distribution

- Do **1,000,000** draws with both Random Walk Metropolis and Gibbs, thinning by 1000

- Do **1,000** draws using Stan's NUTS algorithm (no thinning)

- Do 1,000 independent draws (we can do this for multivariate normal)

Writing a Stan program

What is a Bayesian model?

- Building a Bayesian model forces us to build a model for how the data is generated

- We often think of this as specifying a prior and a likelihood, as if these are two separate things

- They are not!

Generative models

The philosophy behind Stan is to think generatively.

The model is expressed as a joint probability distribution of observed and unobserved variables, which may be decomposed as follows:

$$p(y, \theta) = p(y|\theta)\pi(\theta) \tag{2}$$

The posterior of interest is then proportional to the joint distribution (2):

$$p(\theta|y) \propto p(y|\theta)\pi(\theta) \tag{3}$$

A Bayesian modeller commits to to an a priori joint distribution:

$$p(y, \theta) = \underbrace{p(y|\theta)\pi(\theta)}_{Likelihood \times Prior} = \underbrace{\pi(\theta|y)p(y)}_{Posterior \times Marginal\ Likelihood} \quad (4)$$

Generative models and vague priors

What is the problem with *vague/diffuse* priors?

- If we use an improper prior, then we do not specify a joint model for our data and parameters.

- More importantly, we do not specify a data generating mechanism $p(y)$.

- By construction, these priors do not regularize inferences, which is quite often a bad idea

- Proper but diffuse is better than .improper but is still often problematic.

- If we disallow improper priors, then Bayesian modeling is generative.

- In particular, we have a simple way to simulate from $p(y)$:

$$\theta^\star \sim \pi(\theta)$$
$$\downarrow$$
$$y^\star \sim p(y|\theta^\star)$$

$$y^\star \sim p(y)$$

Stan computations

Stan works in logarithmic terms: all the computations are actually done on log-scale. So, for the posterior we have.

$$\log(\pi(\theta|y)) = \log(\pi(\theta)) + \log(p(y|\theta)) + \mathrm{constant} \tag{5}$$

Products become sums of logs:

$$p(y|\theta) = \prod_{i=1}^{n} p(y_i|\theta) \rightarrow \log(p(y|\theta)) = \sum_{i=1}^{n} \log(p(y_i|\theta)).$$

Starting point

We are now going to write a Stan program together:

- Open a new empty file in RStudio

- Save it as `linear_regression.stan`

Blocks strategy

Stan programs are organized into blocks:

- **data** block: declare data types, sizes, and constraints. Read from data source and constraints validated. Evaluated: once.

- **parameters** block: declare parameter types, sizes, and constraints. Evaluated: every log prob evaluation.

- **transformed parameters** block: declare those parameters transformed from the original ones declared in the parameters block. Evaluated: every log prob evaluation.

- **model** block: statements defining the posterior density in log scale. Evaluated: every log prob evaluation.

- **generated quantities**: declare and define derived variables. (P)RNGs, predictions, event probabilities, decision making. Constraints validated. Evaluated: once per draw.

Data block

```
data {
  // Dimensions
  int<lower=1> N;
  int<lower=1> K;

  // Variables
  matrix[N, K] X;
  vector[N] y;
}
```

```
// single line comment
/* multiple lines of
comments */
```

Parameters' block

```
parameters {
  real alpha;
  vector[K] beta;
  real<lower=0> sigma;
}
```

constraints *required* in
parameters block

```
model {
  // priors  (flat, uniform, if omitted)
  sigma ~ exponential(1);
  alpha ~ normal(0, 10);
  for (k in 1:K) beta[k] ~ normal(0, 5);

  for (n in 1:N) {
    y[n] ~ normal(X[n, ] * beta + alpha, sigma);
  }
}
```

Why is the default automatically uniform?

- $\pi(\theta) \propto 1$ (0 on log scale)
- Nothing added to log prob

Generated quantities block

```stan
generated quantities {
  vector[N] y_rep;
  for (n in 1:N) {
    real y_hat = X[n,] * beta + alpha;   // local/temp
    y_rep[n] = normal_rng(y_hat, sigma);
  }
}
```

```
data {
  int<lower=1>  N;
  int<lower=1>  K;
  matrix[N, K] X;
  vector[N] y;
}
parameters {
  real alpha;
  vector[K] beta;
  real<lower=0> sigma;
}
model {
  sigma ~ exponential(1);
  alpha ~ normal(0, 10);
  for (k in 1:K) beta[k] ~ normal(0, 5);

  for (n in 1:N)
    y[n] ~ normal(alpha + X[n, ] * beta, sigma);
}
generated quantities {
  vector[N] y_rep;
  for (n in 1:N)
    y_rep[n] = normal_rng(alpha + X[n,] * beta, sigma);
}
```

Observed
variables

Unobserved
variables

$$\log \pi(\theta)$$
$$+$$
$$\log \mathrm{p}(\mathbf{y} \mid \theta)$$

Simulate from
generative model

35

Launching the Stan model from R

Now we may launch the Stan program directly in R:

```
library(rstan)

# passing the data (already stored)
data <- list(N=N, K=K, X=X, y=y)

# fitting the model
fit1 <- stan(
  file = 'linear_regression.stan',
  data = data,
  iter = 2000,
  chains = 4)

# extracting the estimates
sims <- extract(fit1)
```

This example studied coaching effects from eight schools.

We denote with y_{ij} the result of the i-th test in the j-th school. We assume the following model:

$$y_{ij} \sim \mathcal{N}(\theta_j, \sigma_y^2)$$
$$\theta_j \sim \mathcal{N}(\mu, \tau^2)$$

Do some schools perform better/worse according to these coaching effects?

Here is the data, already aggregated by schools:

```
schools_dat <- list(J = 8,
                    y = c(28,  8, -3,  7, -1,  1, 18, 12),
                    sigma = c(15, 10, 16, 11,  9, 11, 10, 18))
```

First Stan model: 8 schools

```
// saved as 8schools.stan
data {
  int<lower=0> J;         // number of schools
  real y[J];              // estimated treatment effects
  real<lower=0> sigma[J]; // standard error of effect estimates
}
parameters {
  real mu;                // population treatment effect
  real<lower=0> tau;      // standard deviation in treatment effects
  vector[J] eta;          // unscaled deviation from mu by school
}
transformed parameters {
  vector[J] theta = mu + tau * eta;    // school treatment effects
}
model {
  eta ~ normal(0,1);         // prior
  y ~ normal(theta, sigma);  //likelihood
}
```

First example: 8 schools

To fit the model and visualize the estimates, it is sufficient to type in R
the following commands (with 2000 iterations and 4 chains as a default):

```
fit_8schools <- stan(file = '8schools.stan', data = schools_dat)
print(fit_8schools, pars=c("mu", "tau", "theta"))
```

```
           mean   sd   2.5%   25%    50%    75% 97.5% n_eff Rhat
mu         7.89 5.04  -2.31  4.74   7.92  11.05 18.05  2352    1
tau        6.70 5.71   0.24  2.61   5.43   9.19 21.16  1480    1
theta[1]  11.36 8.23  -2.25  6.18  10.29  15.46 31.15  3161    1
theta[2]   7.89 6.21  -4.43  3.96   7.83  11.78 20.47  4923    1
theta[3]   6.05 7.59 -10.81  1.92   6.56  10.81 20.25  4057    1
theta[4]   7.60 6.44  -5.36  3.74   7.63  11.73 20.57  5055    1
theta[5]   5.13 6.23  -8.45  1.35   5.60   9.26 16.37  4346    1
theta[6]   5.95 6.68  -8.21  1.99   6.30  10.21 18.21  4313    1
theta[7]  10.62 6.93  -1.58  6.12  10.14  14.53 25.63  3381    1
theta[8]   8.40 7.77  -7.18  3.84   8.26  12.63 25.78  3854    1
```

Linked package: `bayesplot`

Posterior graphical analysis with `bayesplot`

Once we fit a model, it is to vital check it via graphical inspection. The
`bayesplot` package (for any help, see the ▸ vignette) is designed to this
task.

The package allows to display:

- Posterior uncertainty intervals

- Univariate marginal posterior distributions

- Bivariate plots

- Trace plots

- Posterior predictive plots

Posterior graphical analysis with `bayesplot`

The first step is to save the posterior. Then you have many choices:

```
library(bayesplot)
posterior <- as.array(fit_8schools)

mcmc_intervals(posterior)      # posterior intervals
mcmc_areas(posterior)          # posterior areas
mcmc_dens(posterior)           # marginal posteriors
mcmc_pairs(posterior)          # bivariate plots
mcmc_trace(posterior)          # trace plots
```

With the arguments pars or regex_pars you may select the desired parameters.
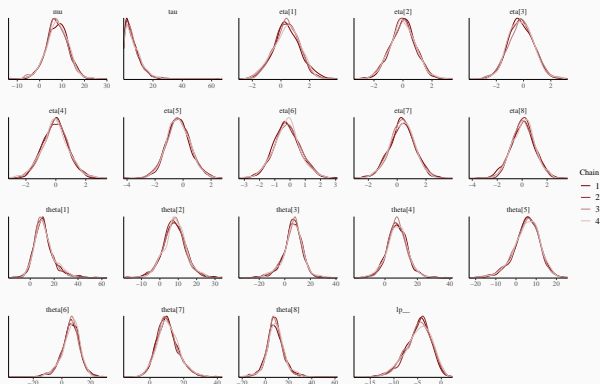
Posterior uncertainty intervals

Posterior intervals

Posterior uncertainty areas

Posterior areas

The Stan shuttle is ready to start! We will learn to:

- **write** simple and more complex model in Stan: lm, glm, hierarchical models.

- **analyze** the posterior summaries.

- **criticize** the model and, eventually, change/reparametrize it.

Application in Stan

Application: HMM for tracking basketball events

https://mc-stan.org/learn-stan/case-studies/bball-hmm.html

Further reading

Further reading:

- Carpenter, B, and Gelman, A, Hoffman, M.D., Lee, D., Goodrich, B.,
  Betancourt, M., Brubaker, M., Guo, J., Li, P., Riddell, A. (2017).
  Stan: A Probabilistic Programming Language, *Journal of statistical
  software 76(1)*. Here the ▶ pdf

Further optional reading about Hamiltonian Monte Carlo:

- Betancourt, M. (2017) A conceptual introduction to Hamiltonian
  Monte Carlo. Here the ▶ pdf